

3. spark.sql

Overview

- Spark SQL is a library included in Spark since version 1.3
- It provides an **easier interface to process tabular data**
- Instead of **RDDs**, we deal with **DataFrames**
- Since Spark 1.6, there is also the concept of **Datasets**, but only for **Scala** and **Java**

SparkContext and SparkSession

- Before Spark 2, there was only SparkContext and SQLContext
- All core functionality was accessed with SparkContext
- All SQL functionality needed the SQLContext, which can be created from an SparkContext
- Since Spark 2, there is now the SparkSession class
- SparkSession is now the **global entry-point** for everything Spark-related

SparkContext and SparkSession

Before Spark 2

```
>>> from pyspark import SparkConf, SparkContext
>>> from pyspark.sql import SQLContext

>>> conf = SparkConf().setAppName(appName).setMaster(master)
>>> sc = SparkContext(conf = conf)
>>> sql_context = new SQLContext(sc)
```

Since Spark 2

```
>>> from pyspark.sql import SparkSession

>>> spark = SparkSession \
    .builder \
    .appName(appName) \
    .master(master) \
    .config("spark.some.config.option", "some-value") \
    .getOrCreate()
```

3.1. DataFrame

DataFrame

- The main entity of Spark SQL is the DataFrame
- A DataFrame is actually an RDD of Rows with a **schema**
- A schema gives the **names of the columns** and their **types**
- Row is a class representing a row of the DataFrame.
- It can be used almost as a **python list**, with its size equal to the number of columns in the schema.

DataFrame

```
>>> from pyspark.sql import Row
```

```
>>> row1 = Row(name="John", age=21)
>>> row2 = Row(name="James", age=32)
>>> row3 = Row(name="Jane", age=18)
>>> row1['name']
'John'
```

```
>>> df = spark.createDataFrame([row1, row2, row3])
>>> df
DataFrame[age: bigint, name: string]
```

```
>>> df.show()
+-----+-----+
| name | age |
+-----+-----+
| John | 21 |
| James | 32 |
| Jane | 18 |
+-----+-----+
```

DataFrame

```
>>> df.printSchema()  
root  
|-- age: long (nullable = true)  
|-- name: string (nullable = true)
```

You can access to the underlying RDD object using `.rdd`

```
>>> print(df.rdd.toDebugString().decode("utf-8"))  
(6) MapPartitionsRDD[16] at javaToPython at NativeMethodAccessorImpl.java:  
| MapPartitionsRDD[15] at javaToPython at NativeMethodAccessorImpl.java:  
| MapPartitionsRDD[14] at javaToPython at NativeMethodAccessorImpl.java:  
| MapPartitionsRDD[13] at applySchemaToPythonRDD at NativeMethodAccessor:  
| MapPartitionsRDD[12] at map at SerDeUtil.scala:137 []  
| MapPartitionsRDD[11] at mapPartitions at SerDeUtil.scala:184 []  
| PythonRDD[10] at RDD at PythonRDD.scala:53 []  
| ParallelCollectionRDD[9] at parallelize at PythonRDD.scala:195 []
```

```
>>> df.rdd.getNumPartitions()  
6
```


Creating DataFrames

- We can use the method `createDataFrame` from the `SparkSession` instance
- Can be used to create a `Spark DataFrame` from:
 - a `pandas.DataFrame` object
 - a local python list
 - an RDD
- Full documentation can be found in the [\[API docs\]](#)

Creating DataFrames

```
>>> rows = [  
    Row(name="John", age=21, gender="male"),  
    Row(name="James", age=25, gender="female"),  
    Row(name="Albert", age=46, gender="male")  
]
```

```
>>> df = spark.createDataFrame(rows)
```

```
>>> df.show()
```

```
+---+-----+-----+  
|age|gender|  name|  
+---+-----+-----+  
| 21|  male|  John|  
| 25|female| James|  
| 46|  male|Albert|  
+---+-----+-----+
```

Creating DataFrames

```
>>> column_names = ["name", "age", "gender"]
>>> rows = [
    ["John", 21, "male"],
    ["James", 25, "female"],
    ["Albert", 46, "male"]
]
>>> df = spark.createDataFrame(rows, column_names)
>>> df.show()
+-----+-----+-----+
|  name|age|gender|
+-----+-----+-----+
|  John| 21|  male|
| James| 25|female|
|Albert| 46|  male|
+-----+-----+-----+
```

Creating DataFrames

```
>>> column_names = ["name", "age", "gender"]
>>> rdd = sc.parallelize([
    ("John", 21, "male"),
    ("James", 25, "female"),
    ("Albert", 46, "male")
])
>>> df = spark.createDataFrame(rdd, column_names)
>>> df.show()
+-----+-----+
|  name|age|gender|
+-----+-----+
|  John| 21|  male|
| James| 25|female|
|Albert| 46|  male|
+-----+-----+
```

3.2. Schemas and types

Schema and Types

- A `DataFrame` always contains a **schema**
- The schema defines the **column names** and **types**
- In all previous examples, the schema was **inferred**
- The schema of a `DataFrame` is represented by the class `types.StructType` [API doc]
- When creating a `DataFrame`, the schema can be either **inferred** or **defined by the user**

```
>>> df.schema  
StructType(List(StructField(name,StringType,true),StructField(age,IntegerType
```

Creating a custom Schema

```
>>> from pyspark.sql.types import *

>>> schema = StructType([
    StructField("name", StringType(), True),
    StructField("age", IntegerType(), True),
    StructField("gender", StringType(), True)
])

>>> rows = [("John", 21, "male")]
>>> df = spark.createDataFrame(rows, schema)
>>> df.printSchema()
>>> df.show()
root
|-- name: string (nullable = true)
|-- age: integer (nullable = true)
|-- gender: string (nullable = true)

+-----+-----+-----+
|name|age|gender|
+-----+-----+-----+
|John| 21|  male|
+-----+-----+-----+
```

Types supported by Spark SQL

- StringType
- IntegerType
- LongType
- FloatType
- DoubleType
- BooleanType
- DateType
- TimestampType
- ...
- The full list of types can be found in [\[API doc\]](#)

3.3. Reading data

Reading data from sources

- Data is usually read from **external sources** (move the **algorithms**, not the **data**)
- Spark SQL provides **connectors** to read from many different sources:
 - Text files (CSV, JSON)
 - Distributed tabular files (Parquet, ORC)
 - In-memory data sources (Apache Arrow)
 - General relational Databases (via JDBC)
 - Third-party connectors to connect to many other databases
 - And you can create your own connector for Spark (in Scala)

Reading data from sources

- In all cases, the syntax is similar:
`spark.read.{source}(path)`
- Spark supports different **file systems** to look for the data:
 - Local files: "`file://path/to/file`" or just "`path/to/file`"
 - HDFS (Hadoop filesystem): "`hdfs://path/to/file`"
 - Amazon S3: "`s3://path/to/file`"

Reading from a CSV file

```
>>> df = spark.read.csv("/path/to/file.csv")
```

```
>>> path = "/path/to/file.csv"  
>>> df = spark.read.option("header", "true").csv(path)
```

```
>>> df = (spark.read  
         .format('csv')  
         .option('header', 'true')  
         .option('sep', ";")  
         .load("/path/to/file.csv")  
         )
```

```
>>> df = spark.read.csv(path, sep=";", header=True)
```

Reading from a CSV file

Main Options

Some important options of the `CSV` reader are listed here:

option	description
<code>sep</code>	The separator character
<code>header</code>	If "true", the first line contains the column names
<code>inferSchema</code>	If "true", the column types will be guessed from the contents
<code>dateFormat</code>	A string representing the format of the date columns

The full list of options can be found in the **API Docs**

Reading from other file types

```
# JSON file
```

```
>>> df = spark.read.json("/path/to/file.json")
```

```
>>> df = spark.read.format("json").load("/path/to/file.json")
```

```
# Parquet file (distributed tabular data)
```

```
>>> df = spark.read.parquet("hdfs://path/to/file.parquet")
```

```
>>> df = spark.read.format("parquet").load("hdfs://path/to/file.parquet")
```

```
# ORC file (distributed tabular data)
```

```
>>> df = spark.read.orc("hdfs://path/to/file.orc")
```

```
>>> df = spark.read.format("orc").load("hdfs://path/to/file.orc")
```

Reading from external databases

- We can use **JDBC** drivers (Java) to read from relational databases
- Examples of databases: **Oracle**, **PostgreSQL**, **MySQL**, etc.
- The **java** driver file must be uploaded to the cluster before trying to access
- This operation can be **very heavy**. When available, specific connectors should be used
- Specific connectors are often provided by **third-party libraries**

Reading from external databases

```
>>> df = spark.read.format("jdbc") \
    .option("url", "jdbc:postgresql:dbserver") \
    .option("dbtable", "schema.tablename") \
    .option("user", "username") \
    .option("password", "p4ssw0rd") \
    .load()
```

or

```
>>> df = spark.read.jdbc(
    url="jdbc:postgresql:dbserver",
    table="schema.tablename"
    properties={
        "user": "username",
        "password": "p4ssw0rd"
    }
)
```


3.4. Queries in `spark.sql`

Performing queries

- Spark SQL was created to be compatible with SQL queries
- So it supports actual SQL queries to be performed on `DataFrames`
- First, the `DataFrame` must be **tagged as a temporary view**
- Then, the queries can be applied using `spark.sql`

Performing queries

```
>>> column_names = ["name", "age", "gender"]
>>> rows = [
    ["John", 21, "male"],
    ["Jane", 25, "female"]
]
>>> df = spark.createDataFrame(rows, column_names)

# Create a temporary view from the DataFrame
>>> df.createOrReplaceTempView("new_view")

# Apply the query
>>> query = "SELECT name, age FROM new_view WHERE gender='male'"
>>> men_df = spark.sql(query)
>>> men_df.show()
+-----+-----+
|name|age|
+-----+-----+
|John| 21|
+-----+-----+
```

Using the API

- Although allowing SQL queries is a very powerful feature, it's **not the best way** to code a complex logic
- Errors are **harder to find** in strings
- Using queries makes the code **less modularizable**
- So Spark SQL provides a **full API** with **SQL-like operations**
- It's **the best way** to code complex logic when using Spark SQL

Basic Operations

operation	description
<code>select</code>	Chooses columns from the table
<code>where</code>	Filters rows based on a boolean rule
<code>limit</code>	Limits the number of rows
<code>orderBy</code>	Sorts the DataFrame based on one or more columns
<code>alias</code>	Changes the name of a column
<code>cast</code>	Changes the type of a column
<code>withColumn</code>	Adds a new column

SELECT

In a SQL query:

```
>>> query = "SELECT name, age FROM table"
```

Using Spark SQL API:

```
>>> df.select("name", "age").show()
```

```
+-----+-----+  
|  name|age|  
+-----+-----+  
|  John| 21|  
|  Jane| 25|  
+-----+-----+
```

WHERE

In a SQL query:

```
>>> query = "SELECT * FROM table WHERE age > 21"
```

Using Spark SQL API:

```
>>> df.where("age > 21").show()
```

Alternatively:

```
>>> df.where(df['age'] > 21).show()
```

```
>>> df.where(df.age > 21).show()
```

```
>>> df.select("*").where("age > 21").show()
```

```
+---+---+---+
|name|age|gender|
+---+---+---+
|Jane| 25|female|
+---+---+---+
```

LIMIT

```
# In a SQL query:  
>>> query = "SELECT * FROM table LIMIT 1"
```

```
# Using Spark SQL API:  
>>> df.limit(1).show()
```

```
# Or even  
>>> df.select("*").limit(1).show()  
+-----+-----+  
|name|age|gender|  
+-----+-----+  
|Jane| 25|female|  
+-----+-----+
```

```
# Note: The result is not deterministic!
```


ORDER BY

In a SQL query:

```
>>> query = "SELECT * FROM table ORDER BY name ASC"
```

Using Spark SQL API:

```
>>> df.orderBy(df.name.asc()).show()
```

```
+-----+-----+-----+
|name|age|gender|
+-----+-----+-----+
|Jane| 25|female|
|John| 21|  male|
+-----+-----+-----+
```

ALIAS (name change)

In a SQL query:

```
>>> query = "SELECT name, age, gender AS sex FROM table"
```

Using Spark SQL API:

```
>>> df.select(df.name, df.age, df.gender.alias('sex')).show()
```

```
+-----+-----+-----+
|name|age|sex|
+-----+-----+-----+
|John|21|male|
|Jane|25|female|
+-----+-----+-----+
```

CAST (type change)

In a SQL query:

```
>>> query = "SELECT name, cast(age AS float) AS age_f FROM table"
```

Using Spark SQL API:

```
>>> df.select(df.name, df.age.cast("float").alias("age_f")).show()
```

Or

```
>>> new_age_col = df.age.cast("float").alias("age_f")
```

```
>>> df.select(df.name, new_age_col).show()
```

```
+-----+-----+
|name|age_f|
+-----+-----+
|John| 21.0|
|Jane| 25.0|
+-----+-----+
```

Adding new columns

In a SQL query:

```
>>> query = "SELECT *, 12*age AS age_months FROM table"
```

Using Spark SQL API:

```
>>> df.withColumn("age_months", df.age * 12).show()
```

Or

```
>>> df.select("*", (df.age * 12).alias("age_months")).show()
```

```
+---+---+-----+-----+
|name|age|gender|(age * 12)|
+---+---+-----+-----+
|John| 21|  male|      252|
|Jane| 25|female|      300|
+---+---+-----+-----+
```

Note: Using withColumn is preferable

Basic operations

- The **full list of operations** that can be applied to a `DataFrame` can be found in the `[DataFrame doc]`
- The **list of operations on columns** can be found in the `[Column docs]`

3.5. Column functions

Column functions

- Oftentimes, we need to make many **transformations using one or more functions**
- Spark SQL has a package called `functions` with many functions available for that
- Some of those functions are only for **aggregations**
Examples: `avg`, `sum`, etc. We will cover them later
- Some others are for **column transformation** or **operations**
Examples: `substr`, `concat`, `datediff`, `floor`, etc.
- The full list is, once again, in the `[API docs]`

Column functions

To use these functions, we first need to import them:

```
>>> from pyspark.sql import functions as fn
```

Note: the "as fn" part is important to **avoid confusion** with native Python functions such as "sum"

Numeric functions examples

```
>>> from pyspark.sql import functions as fn
```

```
>>> columns = ["brand", "cost"]
```

```
>>> df = spark.createDataFrame([
    ("garnier", 3.49),
    ("elseve", 2.71)
], columns)
```

```
>>> round_cost = fn.round(df.cost, 1)
```

```
>>> floor_cost = fn.floor(df.cost)
```

```
>>> ceil_cost = fn.ceil(df.cost)
```

```
>>> df.withColumn('round', round_cost)\
    .withColumn('floor', floor_cost)\
    .withColumn('ceil', ceil_cost)\
    .show()
```

```
+-----+-----+-----+-----+-----+
|  brand|cost|round|floor|ceil|
+-----+-----+-----+-----+-----+
|garnier|3.49|  3.5|    3|    4|
| elseve|2.71|  2.7|    2|    3|
+-----+-----+-----+-----+-----+
```

String functions examples

```
>>> from pyspark.sql import functions as fn

>>> columns = ["first_name", "last_name"]

>>> df = spark.createDataFrame([
    ("John", "Doe"),
    ("Mary", "Jane")
], columns)

>>> last_name_initial = fn.substring(df.last_name, 0, 1)
>>> name = fn.concat_ws(" ", df.first_name, last_name_initial)
>>> df.withColumn("name", name).show()
+-----+-----+-----+
|first_name|last_name|  name|
+-----+-----+-----+
|      John|      Doe|John D|
|      Mary|      Jane|Mary J|
+-----+-----+-----+
```

Date functions examples

```
>>> from datetime import date
>>> from pyspark.sql import functions as fn

>>> df = spark.createDataFrame([
    (date(2015, 1, 1), date(2015, 1, 15)),
    (date(2015, 2, 21), date(2015, 3, 8)),
    ], ["start_date", "end_date"])

>>> days_between = fn.datediff(df.end_date, df.start_date)
>>> start_month = fn.month(df.start_date)

>>> df.withColumn('days_between', days_between)\
    .withColumn('start_month', start_month)\
    .show()
```

start_date	end_date	days_between	start_month
2015-01-01	2015-01-15	14	1
2015-02-21	2015-03-08	15	2

Conditional transformations

- In the `functions` package is a **special function** called `when`
- This function is used to **create a new column** which value **depends on the value of other columns**
- `otherwise` is used to match "the rest"
- Combination between conditions can be done using "&" for "and" and "|" for "or"

Examples

```
>>> df = spark.createDataFrame([
    ("John", 21, "male"),
    ("Jane", 25, "female"),
    ("Albert", 46, "male"),
    ("Brad", 49, "super-hero")
], ["name", "age", "gender"])

>>> supervisor = fn.when(df.gender == 'male', 'Mr. Smith')\
    .when(df.gender == 'female', 'Miss Jones')\
    .otherwise('NA')

>>> df.withColumn("supervisor", supervisor).show()
```

name	age	gender	supervisor
John	21	male	Mr. Smith
Jane	25	female	Miss Jones
Albert	46	male	Mr. Smith
Brad	49	super-hero	NA

User-defined functions

- When you need a **transformation** that is **not available** in the `functions` package, you can create a **User Defined Function** (UDF)
- **Warning:** the performance of this can be **very very low**
- So, it should be used **only** when you are **sure** the operation **cannot be done** with available functions
- To create an UDF, use `functions.udf`, passing a lambda or a named functions
- It is similar to the `map` operation of RDDs

Example

```
>>> from pyspark.sql import functions as fn
>>> from pyspark.sql.types import StringType

>>> df = spark.createDataFrame([(1, 3), (4, 2)], ["first", "second"])

>>> def my_func(col_1, col_2):
    if (col_1 > col_2):
        return "{} is bigger than {}".format(col_1, col_2)
    else:
        return "{} is bigger than {}".format(col_2, col_1)

>>> my_udf = fn.udf(my_func, StringType())

>>> df.withColumn("udf", my_udf(df['first'], df['second'])).show()
+-----+-----+-----+
|first|second|          udf|
+-----+-----+-----+
|    1|    3|3 is bigger than 1|
|    4|    2|4 is bigger than 2|
+-----+-----+-----+
```

3.6. Joins

Performing joins

- Spark SQL supports **joins** between two `DataFrame`
- As in normal SQL, a join **rule** must be defined
- The **rule** can either be a set of **join keys**, or a **conditional rule**
- Join with conditional rules in Spark can be **very heavy**
- **Several types of joins** are available, default is "inner"

Syntax is simple:

```
>>> left_df.join(other=right_df, on=cols, how=join_type)
```

- `cols` contains a column name or a list of column names
- `join_type` is the type of join

Examples

```
>>> from datetime import date
```

```
>>> products = spark.createDataFrame([
    ('1', 'mouse', 'microsoft', 39.99),
    ('2', 'keyboard', 'logitech', 59.99),
], ['prod_id', 'prod_cat', 'prod_brand', 'prod_value'])
```

```
>>> purchases = spark.createDataFrame([
    (date(2017, 11, 1), 2, '1'),
    (date(2017, 11, 2), 1, '1'),
    (date(2017, 11, 5), 1, '2'),
], ['date', 'quantity', 'prod_id'])
```

```
>>> # The default join type is the "INNER" join
```

```
>>> purchases.join(products, 'prod_id').show()
```

prod_id	date	quantity	prod_cat	prod_brand	prod_value
1	2017-11-01	2	mouse	microsoft	39.99
1	2017-11-02	1	mouse	microsoft	39.99
2	2017-11-05	1	keyboard	logitech	59.99

Examples

```
>>> # We can also use a query string (not recommended)
>>> products.createOrReplaceTempView("products")
>>> purchases.createOrReplaceTempView("purchases")
```

```
>>> query = """SELECT * FROM
(purchases AS prc INNER JOIN products AS prd
on prc.prod_id = prd.prod_id)
"""
```

```
>>> spark.sql(query).show()
```

date	quantity	prod_id	prod_id	prod_cat	prod_brand	prod_value
2017-11-01	2	1	1	mouse	microsoft	39.99
2017-11-02	1	1	1	mouse	microsoft	39.99
2017-11-05	1	2	2	keyboard	logitech	59.99

Examples

```
>>> new_purchases = spark.createDataFrame([
    (date(2017, 11, 1), 2, '1'),
    (date(2017, 11, 2), 1, '3'),
], ['date', 'quantity', 'prod_id_x'])
```

```
>>> join_rule = new_purchases.prod_id_x == products.prod_id
```

```
>>> new_purchases.join(products, join_rule, 'left').show()
```

date	quantity	prod_id_x	prod_id	prod_cat	prod_brand	prod_value
2017-11-02	1	3	null	null	null	null
2017-11-01	2	1	1	mouse	microsoft	39.99

Performing joins: some remarks

- Spark **removes the duplicated column** in the DataFrame it outputs after a join operation
- When joining using columns **with nulls**, Spark just skips those

```
>>> df1.show()
+----+-----+
|  id| name|
+----+-----+
| 123|name1|
| 456|name3|
| null|name2|
+----+-----+

>>> df2.show()
+----+-----+
|  id| dept|
+----+-----+
| null|sales|
| 223|Legal|
| 456|  IT|
+----+-----+

>>> df1.join(df2, "id").show
+----+-----+-----+
| id| name| dept|
+----+-----+-----+
|123|name1|sales|
|456|name3|  IT|
+----+-----+-----+
```

Join types

SQL Join Type	In Spark (synonyms)	Description
INNER	"inner"	Data from left and right matching both ways (intersection)
FULL OUTER	"outer", "full", "fullouter"	All rows from left and right with extra data if present (union)
LEFT OUTER	"leftouter", "left"	Rows from left with extra data from right if present
RIGHT OUTER	"rightouter", "right"	Rows from right with extra data from left if present
LEFT SEMI	"leftsemi"	Data from left with a match with right
LEFT ANTI	"leftanti"	Data from left with no match with right
CROSS	"cross"	Cartesian product of left and right (never used)

Join types



Inner join ("inner")

```
>>> inner = df_left.join(df_right, "id", "inner")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

inner

id	value	value
3	A3	A3
4	A4	A4
4	A4	A4_1

Outer join ("outer", "full" or "fullouter")

```
>>> outer = df_left.join(df_right, "id", "outer")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

outer

id	value	value
1	A1	null
2	A2	null
3	A3	A3
4	A4	A4
4	A4	A4_1
5	null	A5
6	null	A6

Left join ("leftouter" or "left")

```
>>> left = df_left.join(df_right, "id", "left")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

left

id	value	value
1	A1	null
2	A2	null
3	A3	A3
4	A4	A4
4	A4	A4_1

Right ("rightouter" or "right")

```
>>> right = df_left.join(df_right, "id", "right")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

right

id	value	value
3	A3	A3
4	A4	A4
4	A4	A4_1
5	null	A5
6	null	A6

Left semi join ("leftsemi")

```
>>> left_semi = df_left.join(df_right, "id", "leftsemi")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

left_semi

id	value
3	A3
4	A4

Left anti joint ("leftanti")

```
>>> left_anti = df_left.join(df_right, "id", "leftanti")
```

df_left

id	value
1	A1
2	A2
3	A3
4	A4

df_right

id	value
3	A3
4	A4_1
4	A4
5	A5
6	A6

left_anti

id	value
1	A1
2	A2

3.7. Aggregations

Performing aggregations

- Maybe **the most used operations** in SQL and Spark SQL
- Similar to SQL, we use "group by" to perform **aggregations**
- We usually can call the aggregation function just after `groupBy`
Namely, we use `groupBy().agg()`
- **Many aggregation functions** in `pyspark.sql.functions`
- Some examples:
 - Numerical: `fn.avg`, `fn.sum`, `fn.min`, `fn.max`, etc.
 - General: `fn.first`, `fn.last`, `fn.count`, `fn.countDistinct`, etc.

Examples

```
>>> from pyspark.sql import functions as fn

>>> products = spark.createDataFrame([
    ('1', 'mouse', 'microsoft', 39.99),
    ('2', 'mouse', 'microsoft', 59.99),
    ('3', 'keyboard', 'microsoft', 59.99),
    ('4', 'keyboard', 'logitech', 59.99),
    ('5', 'mouse', 'logitech', 29.99),
], ['prod_id', 'prod_cat', 'prod_brand', 'prod_value'])

>>> products.groupBy('prod_cat').avg('prod_value').show()

>>> # Or
>>> products.groupBy('prod_cat').agg(fn.avg('prod_value')).show()
+-----+-----+
|prod_cat| avg(prod_value)|
+-----+-----+
|keyboard|          54.99|
|  mouse|43.32333333333333|
+-----+-----+
```


Examples

```
>>> from pyspark.sql import functions as fn

>>> products.groupBy('prod_brand', 'prod_cat')\
    .agg(fn.avg('prod_value')).show()
+-----+-----+-----+
|prod_brand|prod_cat|avg(prod_value)|
+-----+-----+-----+
| microsoft|mouse|49.99|
| logitech|keyboard|49.99|
| microsoft|keyboard|59.99|
| logitech|mouse|29.99|
+-----+-----+-----+
```

Examples

```
>>> from pyspark.sql import functions as fn
```

```
>>> products.groupBy('prod_brand').agg(  
    fn.round(fn.avg('prod_value'), 1).alias('average'),  
    fn.ceil(fn.sum('prod_value')).alias('sum'),  
    fn.min('prod_value').alias('min')  
)
```

```
.show()  
+-----+-----+-----+-----+  
|prod_brand|average|sum|  min|  
+-----+-----+-----+-----+  
|  logitech|    40.0|  80|29.99|  
| microsoft|    53.3|160|39.99|  
+-----+-----+-----+-----+
```

Examples

```
>>> # Using an SQL query
>>> products.createOrReplaceTempView("products")
```

```
>>> query = """
SELECT
  prod_brand,
  round(avg(prod_value), 1) AS average,
  min(prod_value) AS min
FROM products
GROUP BY prod_brand
"""
```

```
>>> spark.sql(query).show()
+-----+-----+-----+
|prod_brand|average|  min|
+-----+-----+-----+
|  logitech|   40.0|29.99|
| microsoft|   53.3|39.99|
+-----+-----+-----+
```

3.8. Window functions

Window functions

- A very, very **powerful feature**
- They allow to solve **complex problems**
- They exist in some relational databases
- A good article about this feature is **[here]**

Window functions

- It's **similar to aggregations**, but the **number of rows doesn't change**
- Instead, **new columns are created**, and the **aggregated values are duplicated** for values of the same "group"
- There are "traditional" aggregations, such as `min`, `max`, `avg`, `sum`
- and "special" types, such as `lag`, `lead`, `rank`

Numerical window functions

```
>>> from pyspark.sql import Window
>>> from pyspark.sql import functions as fn

>>> # First, we create the Window definition
>>> window = Window.partitionBy('prod_brand')

>>> # Then, we can use "over" to aggregate on this window
>>> avg = fn.avg('prod_value').over(window)

>>> # Finally, we can it as a classical column
>>> products.withColumn('avg_brand_value', fn.round(avg, 2)).show()
```

prod_id	prod_cat	prod_brand	prod_value	avg_brand_value
4	keyboard	logitech	49.99	39.99
5	mouse	logitech	29.99	39.99
1	mouse	microsoft	39.99	53.32
2	mouse	microsoft	59.99	53.32
3	keyboard	microsoft	59.99	53.32

Numerical window functions

```
>>> from pyspark.sql import Window
>>> from pyspark.sql import functions as fn

>>> # The window can be defined on multiple columns
>>> window = Window.partitionBy('prod_brand', 'prod_cat')

>>> avg = fn.avg('prod_value').over(window)

>>> products.withColumn('avg_value', fn.round(avg, 2)).show()
```

prod_id	prod_cat	prod_brand	prod_value	avg_value
1	mouse	microsoft	39.99	49.99
2	mouse	microsoft	59.99	49.99
4	keyboard	logitech	49.99	49.99
3	keyboard	microsoft	59.99	59.99
5	mouse	logitech	29.99	29.99

Numerical window functions

```
>>> from pyspark.sql import Window
>>> from pyspark.sql import functions as fn

>>> # Multiple windows can be defined
>>> window1 = Window.partitionBy('prod_brand')
>>> window2 = Window.partitionBy('prod_cat')

>>> avg_brand = fn.avg('prod_value').over(window1)
>>> avg_cat = fn.avg('prod_value').over(window2)

>>> products \
    .withColumn('avg_by_brand', fn.round(avg_brand, 2)) \
    .withColumn('avg_by_cat', fn.round(avg_cat, 2)) \
    .show()
```

prod_id	prod_cat	prod_brand	prod_value	avg_by_brand	avg_by_cat
4	keyboard	logitech	49.99	39.99	54.99
3	keyboard	microsoft	59.99	53.32	54.99
5	mouse	logitech	29.99	39.99	43.32
1	mouse	microsoft	39.99	53.32	43.32
2	mouse	microsoft	59.99	53.32	43.32

Lag and Lead

- `lag` and `lead` are special functions used over an **ordered window**
- They are used to **take the "previous" and "next" value** within the window
- Very useful in datasets with a date column for instance

Lag and Lead

```
>>> purchases = spark.createDataFrame([
    (date(2017, 11, 1), 'mouse'),
    (date(2017, 11, 2), 'mouse'),
    (date(2017, 11, 4), 'keyboard'),
    (date(2017, 11, 6), 'keyboard'),
    (date(2017, 11, 9), 'keyboard'),
    (date(2017, 11, 12), 'mouse'),
    (date(2017, 11, 18), 'keyboard')
], ['date', 'prod_cat'])
```

```
>>> purchases.show()
+-----+-----+
|      date|prod_cat|
+-----+-----+
|2017-11-01|   mouse|
|2017-11-02|   mouse|
|2017-11-04|keyboard|
|2017-11-06|keyboard|
|2017-11-09|keyboard|
|2017-11-12|   mouse|
|2017-11-18|keyboard|
+-----+-----+
```

Lag and Lead

```
>>> window = Window.partitionBy('prod_cat').orderBy('date')
```

```
>>> prev_purch = fn.lag('date', 1).over(window)
```

```
>>> next_purch = fn.lead('date', 1).over(window)
```

```
>>> purchases\  
    .withColumn('prev', prev_purch)\  
    .withColumn('next', next_purch)\  
    .orderBy('prod_cat', 'date')\  
    .show()
```

prod_cat	date	prev	next
keyboard	2017-11-04	null	2017-11-06
keyboard	2017-11-06	2017-11-04	2017-11-09
keyboard	2017-11-09	2017-11-06	2017-11-18
keyboard	2017-11-18	2017-11-09	null
mouse	2017-11-01	null	2017-11-02
mouse	2017-11-02	2017-11-01	2017-11-12
mouse	2017-11-12	2017-11-02	null

Rank, DenseRank and RowNumber

- Another set of **useful "special" functions**
- Also used on **ordered windows**
- They create a **rank**, or an **order** of the items within the window

Rank and RowNumber

```
>>> contestants = spark.createDataFrame([
    ('veterans', 'John', 3000),
    ('veterans', 'Bob', 3200),
    ('veterans', 'Mary', 4000),
    ('young', 'Jane', 4000),
    ('young', 'April', 3100),
    ('young', 'Alice', 3700),
    ('young', 'Micheal', 4000),
], ['category', 'name', 'points'])
```

```
>>> contestants.show()
```

```
+-----+-----+-----+
|category|  name|points|
+-----+-----+-----+
|veterans|  John| 3000|
|veterans|   Bob| 3200|
|veterans|  Mary| 4000|
|  young|  Jane| 4000|
|  young| April| 3100|
|  young| Alice| 3700|
|  young|Micheal| 4000|
+-----+-----+-----+
```

Rank and RowNumber

```
>>> window = Window.partitionBy('category')\
    .orderBy(contestants.points.desc())

>>> rank = fn.rank().over(window)
>>> dense_rank = fn.dense_rank().over(window)
>>> row_number = fn.row_number().over(window)

>>> contestants\
    .withColumn('rank', rank)\
    .withColumn('dense_rank', dense_rank)\
    .withColumn('row_number', row_number)\
    .orderBy('category', fn.col('points').desc())\
    .show()
```

category	name	points	rank	dense_rank	row_number
veterans	Mary	4000	1	1	1
veterans	Bob	3200	2	2	2
veterans	John	3000	3	3	3
young	Jane	4000	1	1	1
young	Micheal	4000	1	1	2
young	Alice	3700	3	2	3
young	April	3100	4	3	4

3.9. Writing dataframes

Writing dataframes

- Very **similar to reading**. Output formats are the same: **csv**, **json**, **parquet**, **orc**, **jdbc**, etc. Note that **write is an action**
- Instead of `df.read.{source}` use `df.write.{target}`
- Main option is **mode** with possible values:
 - **"append"**: append contents of this **DataFrame** to existing data.
 - **"overwrite"**: overwrite existing data
 - **"error"**: throw an exception if data already exists
 - **"ignore"**: silently ignore this operation if data already exists.

Example

```
>>> products.write.csv('/products.csv')
>>> products.write.mode('overwrite').parquet('/file.parquet')
>>> products.write.format('parquet').save('/file.parquet')
```

3.10. Under the hood...


Query planning and optimization

A **lot happens under the hood** when executing an **action** on a `DataFrame`. The query goes through the following **execution stages**:

1. Logical Analysis
2. Caching Replacement
3. Logical Query Optimization (using rule-based and cost-based optimizations)
4. Physical Planning
5. Physical Optimization (e.g. Whole-Stage Java Code Generation or Adaptive Query Execution)
6. Constructing the RDD of Internal Binary Rows (that represents the structured query in terms of Spark Core's RDD API)

<https://jaceklaskowski.gitbooks.io/mastering-spark-sql/spark-sql.html>

Query planning and optimization



<https://jaceklaskowski.gitbooks.io/mastering-spark-sql/spark-sql.html>